



专业云计算服务商

深度学习TensorFlow下的计算机视觉

<http://hacker.duanshishi.com>

<https://github.com/burness>

Burness Duan

UCloud



专业云计算服务商

OUTLINE

1. TensorFlow Features
2. TensorFlow And Deep CV
3. TFLearn Introduction



专业云计算服务商

TensorFlow Features

1. Programming Model and Basic Concepts
2. Implementation
3. Extensions
4. Optimizations

Abadi M, Agarwal A, Barham P, et al. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015[J]. Software available from tensorflow.org, 2015, 1.



专业云计算服务商

Programming Model and Basic Concepts

```
import tensorflow as tf

b = tf.Variable(tf.zeros([100])) # 100-d vector, init to zeroes
W = tf.Variable(tf.random_uniform([784,100],-1,1)) # 784x100 matrix w/rnd vals
x = tf.placeholder(name="x") # Placeholder for input
relu = tf.nn.relu(tf.matmul(W, x) + b) # Relu(Wx+b)
C = [...] # Cost computed as a function # of Relu

s = tf.Session()
for step in xrange(0, 10):
    input = ...construct 100-D input array ... # Create 100-d vector for input
    result = s.run(C, feed_dict={x: input}) # Fetch cost, feeding x=input
    print step, result
```

Figure 1: Example TensorFlow code fragment




Figure 2: Corresponding computation graph for Figure 1



专业云计算服务商

Programming Model and Basic Concepts

1. Operations and Kernels

2. Sessions

3. Variables



Programming Model and Basic Concepts

Operation: An abstract computation, can have attributes

Kernel: A particular implementation of an operation that can be run on a particular type of device (CPU or GPU)

Sessions: Interact with the TensorFlow system, Run the full graph or a few distinct subgraphs

Variables: A special operation returns a handle to a persistent mutable tensor

https://www.tensorflow.org/versions/r0.11/how_tos/adding_an_op/index.html



专业云计算服务商

Implementation

1. Single-Device Execution

2. Multi-Device Execution

- a) Node Placement
- b) Cross-Device Communication

3. Distributed Execution



Single-Device Execution

The nodes of the graph are executed in an order that respects the dependencies between the nodes



专业云计算服务商

Multi-Device Execution


Node Placement: Based on heuristics associated with different operation types or is measured based on an actual set of placement decisions for earlier extensions of the graph



专业云计算服务商

Multi-Device Execution

Cross-Device Communication:





专业云计算服务商

Distributed Execution

Similar to multi-device execution, After device placement, a subgraph is created per device. Send/Receive node communicate across worker processes use remote communication such as TCP or RDMA.



专业云计算服务商

Distributed TensorFlow Example

[Distributed_tensorflow.py](#)

```
# On ps0.example.com:
$ python trainer.py \
    --ps_hosts=ps0.example.com:2222,ps1.example.com:2222 \
    --worker_hosts=worker0.example.com:2222,worker1.example.com:2222 \
    --job_name=ps --task_index=0
# On ps1.example.com:
$ python trainer.py \
    --ps_hosts=ps0.example.com:2222,ps1.example.com:2222 \
    --worker_hosts=worker0.example.com:2222,worker1.example.com:2222 \
    --job_name=ps --task_index=1
# On worker0.example.com:
$ python trainer.py \
    --ps_hosts=ps0.example.com:2222,ps1.example.com:2222 \
    --worker_hosts=worker0.example.com:2222,worker1.example.com:2222 \
    --job_name=worker --task_index=0
# On worker1.example.com:
$ python trainer.py \
    --ps_hosts=ps0.example.com:2222,ps1.example.com:2222 \
    --worker_hosts=worker0.example.com:2222,worker1.example.com:2222 \
    --job_name=worker --task_index=1
```


Extensions

1. Gradient Computation ↔ TF hold a lot of GPU memory
2. Partial Execution ↔ Run the nodes with the deps of the out tensor
3. Device Constraints ↔ You can define which nodes run on the specify devices
4. Control Flow ↔ Allow us to skip the execution of an subgraph based on the Value of a boolean tensor
5. Input Operations ↔ From the storage system to the worker
6. Queues ↔ Allow the inpu data to be prefetched from disk files while a previous batch of data is being processed
7. Containers ↔ Share state even across completely disjoint computation graphs



专业云计算服务商


Gradient Computation





专业云计算服务商

Partial Execution





Device Constraints

Client can control the placement of nodes on devices by providing partial constraints for a node about Which devices it can execute on.



专业云计算服务商

Control Flow

High-level programming constructs such if-conditionals and while-loops can be easily compiled into dataflow graph with these control flow operators.



Input Operations

An special node in graph: efficient mechanism used for training large-scale machine learning models



Queues

Allow input data to be prefetched from disk files while a previous batch of data is still being processed by the model



Containers

Allow input data to be prefetched from disk files while a previous batch of data is still being processed by the model



Optimizations

1. Common Subexpression Elimination
2. Asynchronous Kernels
3. Optimized Libraries for Kernel Implementation



Common Subexpression Elimination

Convert the redundant copies of the same computation to just a single one



Asynchronous Kernels

Non-blocking kernels, avoid tying up an execution thread for unbounded periods of time while waiting for I/O or other events.



Optimized Libs For Kernel Implementations

Eigen\BLAS\cuBLAS, GPU CNN Kernel: cuda-convnet, cuDNN



专业云计算服务商

TensorFlow And Deep CV

1. Image Classification

2. Neural Style

3. Txt2Img, img2txt



专业云计算服务商

Image Classification


1. Training from scratch
2. Retrain from pre-trained model
3. Load model and frozen some layers' weights and retrain the other layers



专业云计算服务商

Training from scratch

```
mnist = input_data.read_data_sets(FLAGS.data_dir, one_hot=True)
x = tf.placeholder(tf.float32, [None, 784])
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))
y = tf.matmul(x, W) + b
y_ = tf.placeholder(tf.float32, [None, 10])
cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y, y_))
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
sess = tf.InteractiveSession()
tf.initialize_all_variables().run()
for _ in range(1000):
    batch_xs, batch_ys = mnist.train.next_batch(100)
    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
    correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print(sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels}))
```






专业云计算服务商

Retrain from pre-trained model

1. Load pre params from a model file(pb, ckpt)
2. Change your model to the new task(class num)
3. Fit the train data to update the all weights
4. Fit your test data to inference






专业云计算服务商

Retrain from pre-trained model (frozen)

1. Load pre params from a model file(pb)
2. Frozen your layers' weight
3. Change the network's class num
4. Fit your train data to update the unfrozen layers' weight
5. Fit your test data to inference



https://github.com/spark-mler/WorkWithTensorflow/tree/master/cv_bot/models/pretrain_inference



专业云计算服务商

Neural Style

1.MRF-Based

2.CNN-Based

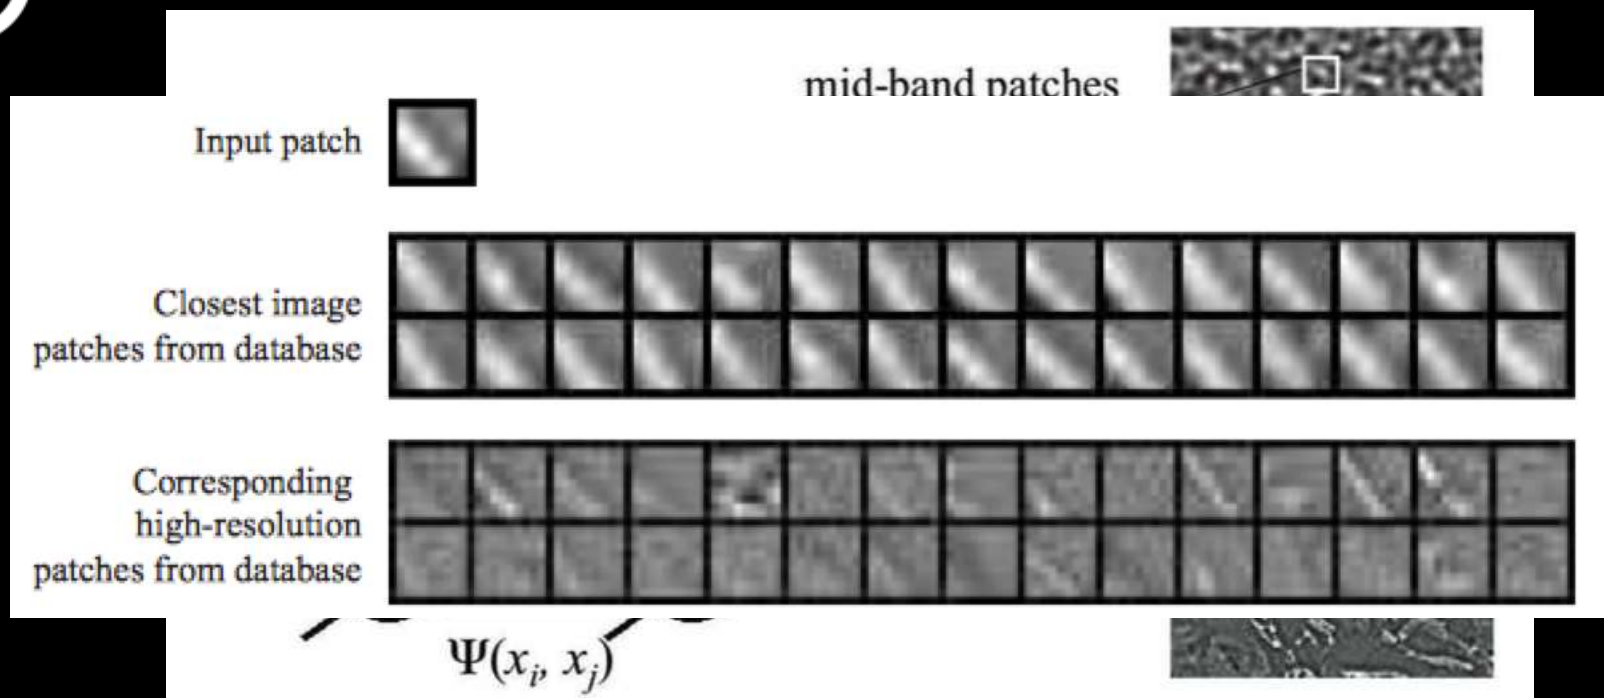
3.MRF and CNN-Based

4.Fast Neural Style



专业云计算服务商

MRF-Based



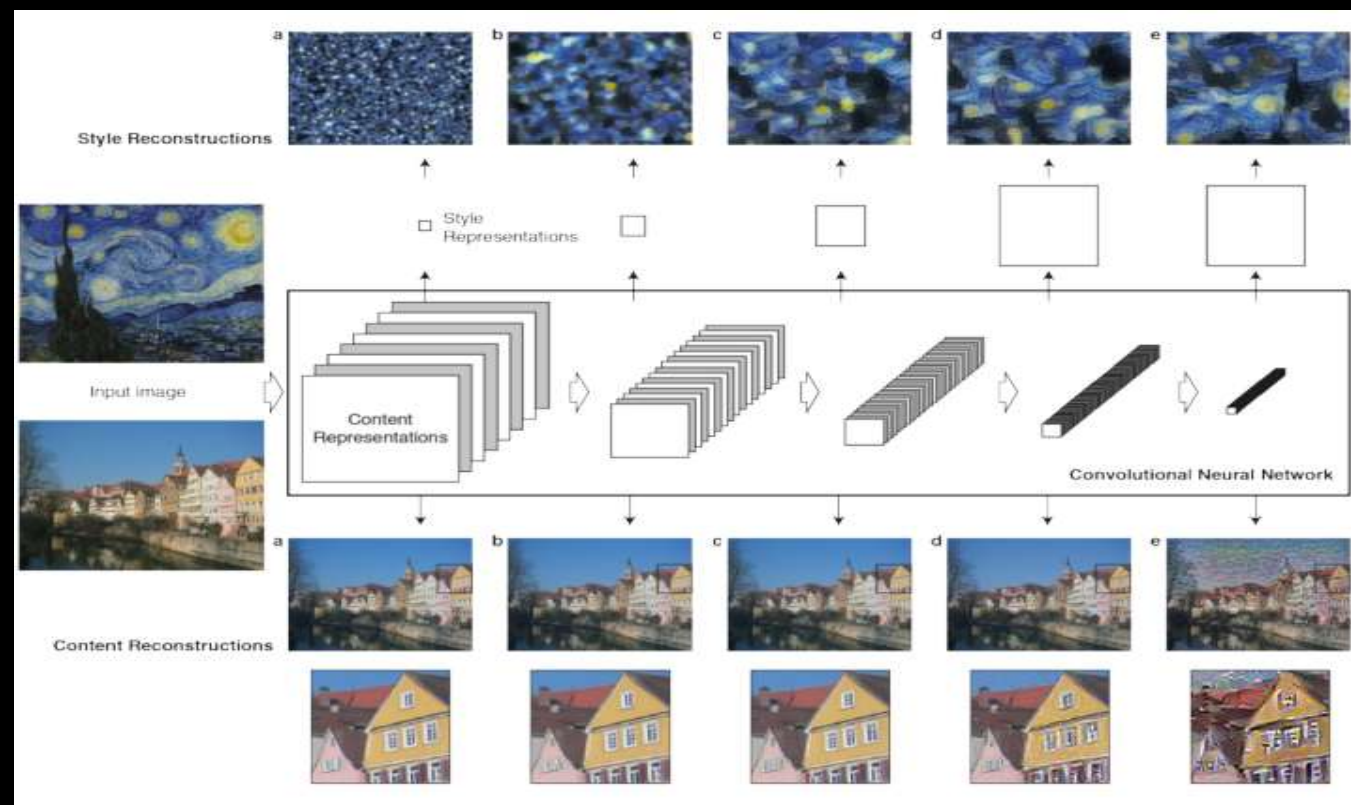
Freeman W T, Liu C. Markov random fields for super-resolution and texture synthesis[J]. Advances in Markov Random Fields for Vision and Image Processing, 2011, 1: 155-165.

Efros A A, Leung T K. Texture synthesis by non-parametric sampling[C]//Computer Vision, 1999. The Proceedings of the Seventh IEEE International Conference on. IEEE, 1999, 2: 1033-1038.



专业云计算服务商

CNN-Based



Gatys L A, Ecker A S, Bethge M. A neural algorithm of artistic style[J]. arXiv preprint arXiv:1508.06576, 2015.

<https://github.com/anishathalve/neural-style>

<https://github.com/jcjohnson/neural-style>



MRF And CNN-Based



Input style



Input content



Gatys et al



Ours

$$\mathbf{x} = \arg \min_x E_s(\Phi(\mathbf{x}), \Phi(\mathbf{x}_s)) + \alpha_1 E_c(\Phi(\mathbf{x}), \Phi(\mathbf{x}_c)) + \alpha_2 \Upsilon(\mathbf{x})$$


Li C, Wand M. Combining Markov Random Fields and Convolutional Neural Networks for Image Synthesis[J]. arXiv preprint arXiv:1601.04589, 2016

<https://github.com/chuanli11/CNNMRF>



专业云计算服务商

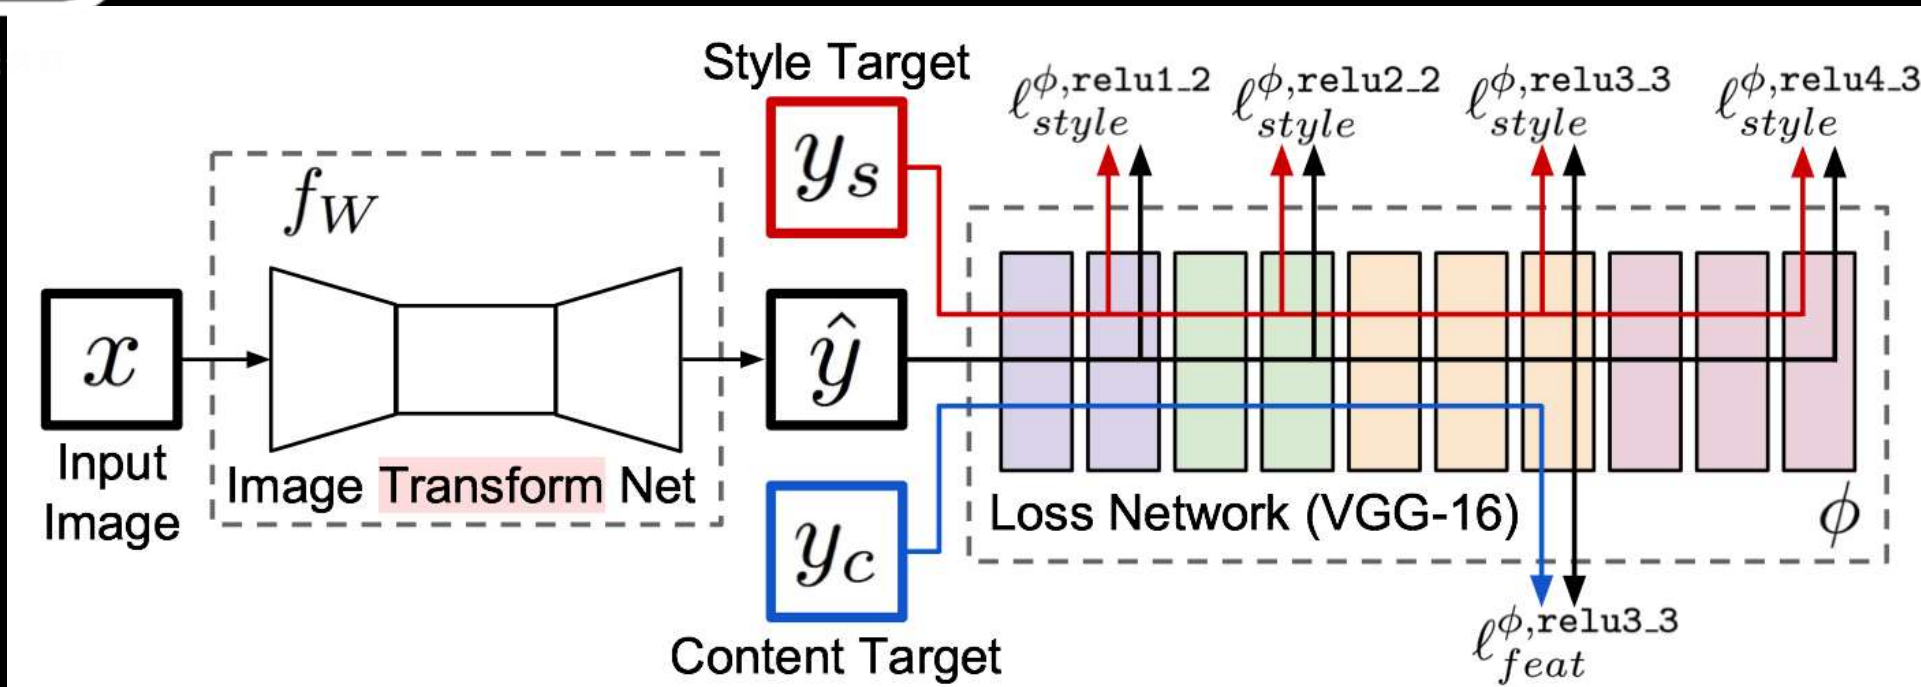
MRF And CNN-Based





专业云计算服

Fast Neural Style




Johnson J, Alahi A, Fei-Fei L. Perceptual losses for real-time style transfer and super-resolution[J]. arXiv preprint arXiv:1603.08155, 2016.

https://github.com/burness/neural_style_tensorflow/tree/master/fast_neural_style



专业云计算服务商

Fast Neural Style





专业云计算服务商

Text↔Image

1.Text-to-Image

2.Image-to-Text



Text-to-Image

Generative Adversarial Network

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Generator G and a discriminator D compete in a two-player minimax game

G: fool the discriminator

D: Distinguish real training data from synthetic images

Reed S, Akata Z, Yan X, et al. Generative adversarial text to image synthesis[J]. arXiv preprint arXiv:1605.05396, 2016.

<https://github.com/paarthneekhara/text-to-image>

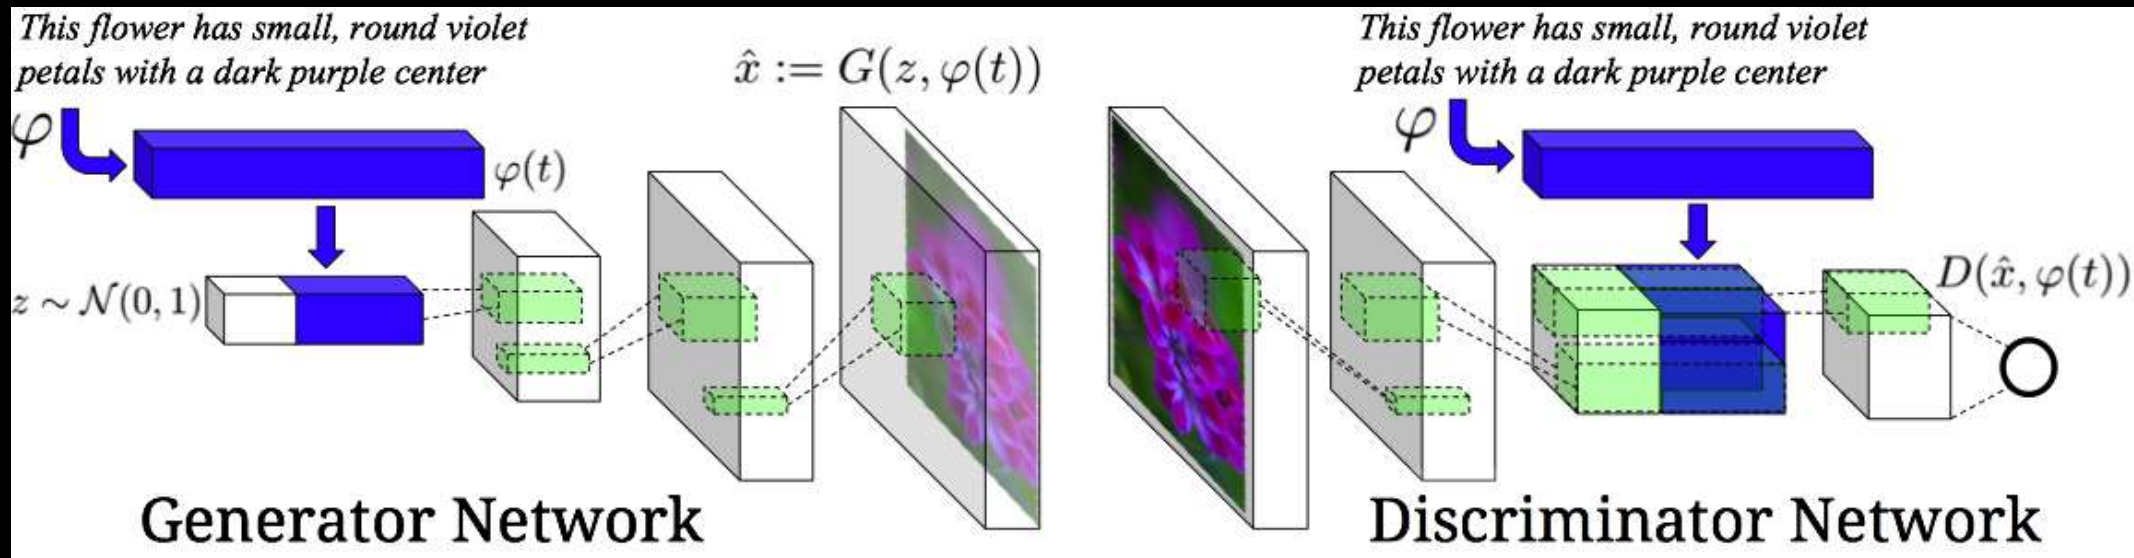
<http://www.jianshu.com/p/db87c51de510>



专业云计算服务商

Text-to-Image

Generative Adversarial Network



$$G: \mathbb{R}^Z \times \mathbb{R}^T \rightarrow \mathbb{R}^D$$

$$D: \mathbb{R}^T \times \mathbb{R}^D \rightarrow \{0,1\}$$

Text-to-Image

Results

the flower shown has yellow anther red pistil and bright red petals



this flower has petals that are yellow, white and purple and has dark lines



the petals on this flower are white with a yellow center




this flower has a lot of small round pink petals



this flower is orange in color, and has petals that are ruffled and rounded



the flower has yellow petals and the center of it is brown





专业云计算服务商

Image-to-Text

A person on a beach flying a kite.



A black and white photo of a train on a train track.



A person skiing down a snow covered slope.



A group of giraffe standing next to each other.



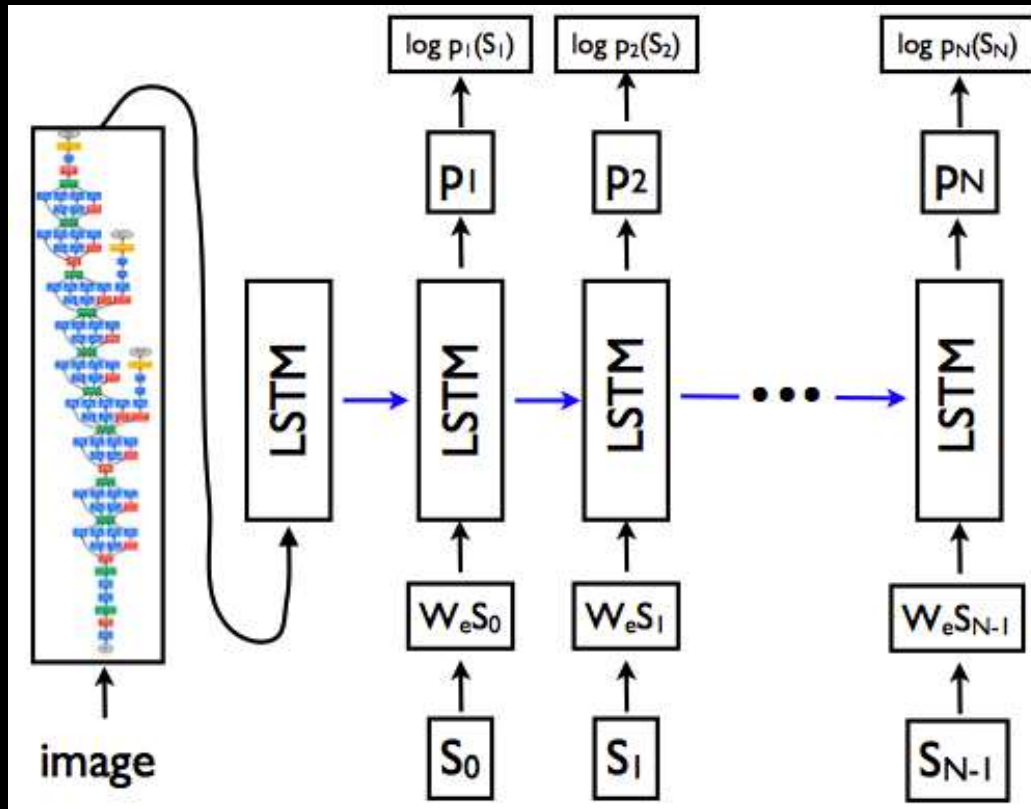
Vinyals O, Toshev A, Bengio S, et al. Show and Tell: Lessons learned from the 2015 MSCOCO Image Captioning Challenge[J]. 2016.

<https://github.com/tensorflow/models/tree/master/im2txt>

<https://github.com/tensorflow/models/issues/480>

<https://github.com/tensorflow/models/pull/485/commits/c6a4f783080c5310ce0e3244daa31af57df12def>

Image-to-Text



$$\theta^* = \arg \max_{\theta} \sum_{(I, S)} \log p(S|I; \theta)$$

$$\log p(S|I) = \sum_{t=0}^N \log p(S_t|I, S_0, \dots, S_{t-1})$$

$$\begin{aligned} x_{-1} &= \text{CNN}(I) \\ x_t &= W_e S_t, \quad t \in \{0 \dots N-1\} \\ p_{t+1} &= \text{LSTM}(x_t), \quad t \in \{0 \dots N-1\} \end{aligned}$$

$$L(I, S) = - \sum_{t=1}^N \log p_t(S_t)$$

Ioffe S, Szegedy C. Batch normalization: Accelerating deep network training by reducing internal covariate shift[J]. arXiv preprint arXiv:1502.03167, 2015.

Deep Visual-Semantic Alignments for Generating Image Descriptions

Andrej Karpathy, Li Fei-Fei



微信号: transwarp-sh

Captions for image 54.jpg:

- 0) a group of people standing next to each other . (p=0.002586)
- 1) a group of people standing in a room . (p=0.000496)
- 2) a group of people standing next to each other in a room . (p=0.000119)

amax@am tensorflow/models/im2txt\$



Captions for image nba_jandunks_01.jpg:


- 0) a group of young men playing a game of basketball . (p=0.005181)
- 1) a group of men playing a game of basketball . (p=0.004607)
- 2) a group of men playing basketball on a court (p=0.000636)



专业云计算服务商

TFlearn Introduction

```
with tf.name_scope('conv1'):
    W = tf.Variable(tf.random_normal([5, 5, 1, 32]), dtype=tf.float32, name='Weights')
    b = tf.Variable(tf.random_normal([32]), dtype=tf.float32, name='biases')
    x = tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')
    x = tf.add_bias(W, b)
    x = tf.nn.relu(x)
```



```
tflearn.conv_2d(x, 32, 5, activation='relu', name='conv1')
```


TFLearn Example

```
from __future__ import division, print_function, absolute_import

import tflearn
from tflearn.layers.core import input_data, dropout, fully_connected
from tflearn.layers.conv import conv_2d, max_pool_2d
from tflearn.layers.normalization import local_response_normalization
from tflearn.layers.estimator import regression


import tflearn.datasets.oxflower17 as oxflower17
X, Y = oxflower17.load_data(one_hot=True, resize_images=True)

# Building 'AlexNet'
network = input_data(shape=[None, 227, 227, 3])
network = conv_2d(network, 96, 11, strides=4, activation='relu', batch_normalization=True)
network = max_pool_2d(network, 3, strides=2)
network = local_response_normalization(network)
network = conv_2d(network, 256, 5, activation='relu', batch_normalization=True)
network = max_pool_2d(network, 3, strides=2)
network = local_response_normalization(network)
network = conv_2d(network, 384, 3, activation='relu', batch_normalization=True)
network = conv_2d(network, 384, 3, activation='relu', batch_normalization=True)
network = conv_2d(network, 256, 3, activation='relu', batch_normalization=True)
network = max_pool_2d(network, 3, strides=2)
network = local_response_normalization(network)
network = fully_connected(network, 4096, activation='relu')
network = dropout(network, 0.5)
network = fully_connected(network, 4096, activation='relu')
network = dropout(network, 0.5)
network = fully_connected(network, 17, activation='softmax')
network = regression(network, optimizer='momentum',
                      loss='categorical_crossentropy',
                      learning_rate=0.001)


# Training
model = tflearn.DNN(network, checkpoint_path='model_alexnet',
                    max_checkpoints=1, tensorboard_verbose=2)
model.fit(X, Y, n_epoch=1000, validation_set=0.1, shuffle=True,
        show_metric=True, batch_size=64, snapshot_step=200,
        snapshot_epoch=False, run_id='alexnet_oxflowers17')
```

- Accuracy


- Accuracy/



- Accuracy/ (raw)




- Accuracy/Validation




- Loss


- Loss/



- Loss/ (raw)



- Loss/Validation



Run	Value	Step	Time	Relative
	0.02387	2.923k	Sat Jul 16, 23:18:22	3h 6m 20s





专业云计算服务商

THANKS
&
QA